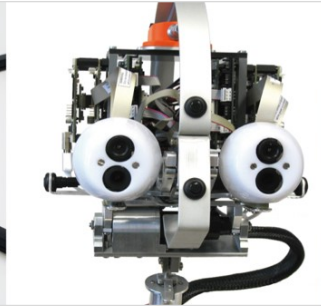
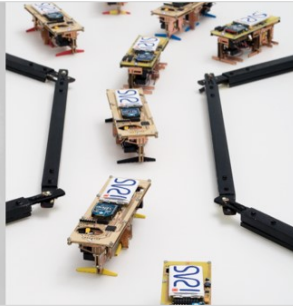
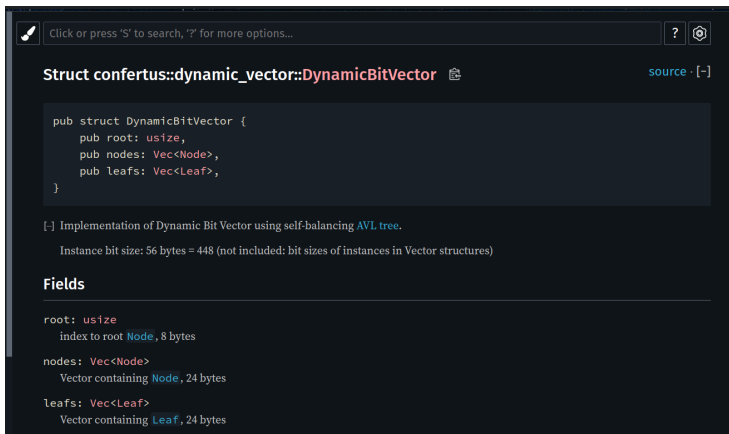


confertus: a dynamic bitvector implementation in rust


Felix Karg | 25. Juli 2022



Automatically Generated documentation



Click or press 'S' to search, '?' for more options...

Struct confertus::dynamic_vector::DynamicBitVector  [source](#) · [-]

```
pub struct DynamicBitVector {
    pub root: usize,
    pub nodes: Vec<Node>,
    pub leaves: Vec<Leaf>,
}
```

[-] Implementation of Dynamic Bit Vector using self-balancing [AVL tree](#).

Instance bit size: 56 bytes = 448 (not included: bit sizes of instances in Vector structures)

Fields

root: `usize`
index to root [Node](#), 8 bytes

nodes: `Vec<Node>`
Vector containing [Node](#), 24 bytes

leaves: `Vec<Leaf>`
Vector containing [Leaf](#), 24 bytes

Documentation can be found at <https://fkarg.me/confertus>

Documentation

Structure Visualization

Code Examples

LOC-Stats

Results

●

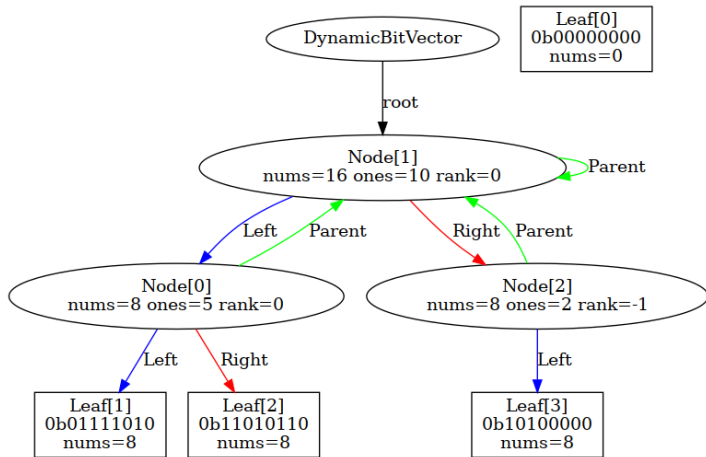
○

○○○○○○○

○

○○○

Structure Visualization with dotviz



Documentation

○

Structure Visualization

●

Code Examples

○○○○○○

LOC-Stats

○

Results

○○○

DynamicBitVector

```
+-----+
|DynamicBitVector |
+-----+
|root: usize      |
|nodes: Vec<Node> |
|leafs: Vec<Leaf> |
+-----+
```

```
1  /// Implementation of Dynamic Bit Vector based on
2  /// a self-balancing [AVL tree]. ...
3  #[derive(Debug, PartialEq, Clone, Default)]
4  pub struct DynamicBitVector {
5      /// index to root [`Node`]
6      pub root: usize, // 8 bytes
7      // positively indexed, usize
8      /// Vector containing [`Node`]
9      pub nodes: Vec<Node>, // 24 bytes
10     // negatively indexed, isize
11     /// Vector containing [`Leaf`]
12     pub leafs: Vec<Leaf>, // 24 bytes
13 }
```

Documentation

○

Structure Visualization

○

Code Examples

●○○○○○

LOC-Stats

○

Results

○○○

Node

```
+-----+
|Node      |
+-----+
|parent: Option<usize>|
|left: Option<isize>  |
|right: Option<isize> |
|nums: usize          |
|ones: usize          |
|rank: i8             |
+-----+
```

```
1 pub struct Node {
2     /// index of parent Node
3     pub parent: Option<usize>, // 8 bytes + 1bit
4     /// left side subtree child index
5     pub left: Option<isize>, // 8 bytes + 1bit
6     /// right side subtree child index
7     pub right: Option<isize>, // 8 bytes + 1bit
8     /// number of filled bits on the left subtree
9     pub nums: usize, // 8 bytes
10    /// number of ones on the left subtree
11    pub ones: usize, // 8 bytes
12    /// difference of height between left and
13    /// right subtree. Valid values are -1, 0, 1
14    pub rank: i8, // 1 byte
15 }
```

Documentation



Structure Visualization



Code Examples



LOC-Stats



Results



Leaf

```
+-----+
| Leaf   |
+-----+
| parent: usize |
| value: u128   |
| nums: u8      |
+-----+
```

```
1  /// Leaf element of [`DynamicBitVector`]
2  /// ...
3  #[derive(PartialEq, Clone, Default)]
4  pub struct Leaf {
5      /// reference to parent [`Node`]
6      pub parent: usize, // 8 bytes
7      /// container for actual bit values
8      pub value: u128, // 16 bytes
9      /// number of bits used in `value`-container
10     pub nums: u8, // 1 byte
11     // realistically below u128::BITS,
12     // so u8::MAX = 255 is sufficient
13 }
```

Implementation of rank and select for u64

```
1 use core::arch::x86_64::{_pdep_u64, _popcnt64, _tzcnt_u64};
2
3 #[inline]
4 #[cfg(all(target_arch = "x86_64", target_feature = "bmi1", target_feature = "bmi2"))]
5 unsafe fn select_internal(&self, bit: bool, n: usize) -> usize {
6     _tzcnt_u64(_pdep_u64(1 << n, if bit { *self } else { !self })) as usize
7 }
8
9 #[inline]
10 #[cfg(target_arch = "x86_64")]
11 unsafe fn rank_internal(&self, bit: bool, index: usize) -> usize {
12     _popcnt64({if bit {*self} else {!self}}.overflowing_shl(u64::BITS - index as u32).0 as i64)
13         as usize
14 }
```

(I also have respective non-architecture-specific fallbacks for both)

Documentation



Structure Visualization



Code Examples



LOC-Stats



Results



Index for DynamicBitVector

I have **two** implementations for indexing!

```
1 fn index(&DynamicBitVector self, index: usize) -> Node { ... }  
2 fn index(&DynamicBitVector self, index: isize) -> Leaf { ... }
```

This way, I can do type-safe indexing:

- self[+3 as usize] -> Node
- self[+3 as isize] -> Leaf
- self[-3 as isize] -> Leaf

Unified Traversal: Apply

```
1 pub fn apply<T>(&mut self, mut f: FnMut(...) -> T, index: usize) -> T {
2     self.apply_node(self.root, f, index)
3 }
4 fn apply_node<T>(&mut self, node: usize, mut f: FnMut(...) -> T, index: usize) -> T {
5     if self[node].nums <= index {
6         let right_id = self[node].right.unwrap();
7         if right_id >= 0 {
8             self.apply_node(right_id as usize, f, index - self[node].nums)
9         } else {
10             f(self, right_id, index - self[node].nums)
11         }
12     } else {
13         let left_id = self[node].left.unwrap();
14         if left_id >= 0 {
15             self.apply_node(left_id as usize, f, index)
16         } else {
17             f(self, left_id, index)
18         }
19     }
20 }
```

Documentation

Structure Visualization

Code Examples

LOC-Stats

Results



Beispielimplementation: flip

```
1  impl DynamicBitVector {  
2      ...  
3  
4      #[inline]  
5      pub fn flip(&mut self, index: usize) {  
6          let leaf = self.apply(Self::flip_leaf, index);  
7          self.update_left_values(self[leaf].parent, leaf);  
8      }  
9  
10     #[inline]  
11     fn flip_leaf(&mut self, leaf: isize, index: usize) -> isize {  
12         self[leaf].flip(index);  
13         leaf  
14     }  
15 }
```

Documentation

○

Structure Visualization

○

Code Examples

○○○○○○●

LOC-Stats

○

Results

○○○

LOC-Stats

\$ tokei

Language	Files	Lines	Code	Comments	Blanks
Rust	18	3510	2740	389	381
- Markdown	15	434	0	389	45
(Total)		3944	2740	778	426

Documentation



Structure Visualization



Code Examples



LOC-Stats



Results



Time invested



Does not include time needed to create this presentation

Documentation

○

Structure Visualization

○

Code Examples

○○○○○○○

LOC-Stats

○

Results

●○○

Results

Unfortunately, I didn't fully get it to run until the end.

```
$ RUSTFLAGS="-C target-cpu=native" cargo run --release bv input/example_bv_10k.txt out.txt
  Compiling confertus v0.1.0 (/home/pars/Coding/confertus)
  Finished release [optimized] target(s) in 4.11s
  Running `target/release/confertus bv input/example_bv_10k.txt out.txt`
index 7 out of bounds for 5
Error: "Leaf.insert: Index out of bounds `index > self.nums`"
```

Turns out, it doesn't uphold all invariances regarding nums and ones, causing them to drift – which eventually causes frustrating errors like this.

What are your Questions?

Documentation



Structure Visualization



Code Examples



LOC-Stats



Results



Abstraction: Static Bit Vector

```
1  /// Functions associated with static bit vectors. ...
2  pub trait StaticBitVec {
3      ...
4
5      /// Return number of on-bits in container or on left half for tree-based elements
6      fn ones(&self) -> usize;
7
8      /// Access bit value at position `index`
9      fn access(&self, index: usize) -> bool;
10
11     /// Returns number of `bit`-values up to `index` in container
12     fn rank(&self, bit: bool, index: usize) -> usize;
13
14     /// Return index of `n`-th `bit`-value in container
15     fn select(&self, bit: bool, n: usize) -> usize;
16 }
```

Abstractions



Abstraction: Dynamic Bit Vector

```
1  /// Functions associated with dynamic bit vectors.
2  pub trait DynBitVec: StaticBitVec {
3      /// Insert `bit` at position `index` in underlying container
4      fn insert(&mut self, index: usize, bit: bool) -> Result<(), &'static str>;
5
6      /// Remove bit value at position `index`
7      fn delete(&mut self, index: usize) -> Result<(), &'static str>;
8
9      /// Flip bit at position `index`, updates `ones` and `num` values accordingly
10     fn flip(&mut self, index: usize);
11
12     /// Return used capacity of underlying container
13     fn nums(&self) -> usize;
14 }
```

Abstractions

